[D.3.2] Universal AI: AIXI

David Quarel

Iliad Intensive

June 2026

The big picture: one agent, two halves

- **Universal AI** asks: what would an *optimally intelligent* agent look like, given unlimited compute and the weakest possible assumptions?

The big picture: one agent, two halves

- **Universal AI** asks: what would an *optimally intelligent* agent look like, given unlimited compute and the weakest possible assumptions?
- It splits into two problems:
  1. **Prediction (Solomonoff):** passively observe a sequence $x_1, x_2, \ldots$ and predict the next symbol. *How should you learn?*

The big picture: one agent, two halves

- **Universal AI** asks: what would an *optimally intelligent* agent look like, given unlimited compute and the weakest possible assumptions?
- It splits into two problems:
  1. **Prediction (Solomonoff):** passively observe a sequence $x_1, x_2, \ldots$ and predict the next symbol. *How should you learn?*
  2. **Action (AIXI):** interact with an unknown environment, take actions, collect reward. *How should you act?*

The big picture: one agent, two halves

- **Universal AI** asks: what would an *optimally intelligent* agent look like, given unlimited compute and the weakest possible assumptions?
- It splits into two problems:
  1. **Prediction (Solomonoff):** passively observe a sequence $x_1, x_2, \ldots$ and predict the next symbol. *How should you learn?*
  2. **Action (AIXI):** interact with an unknown environment, take actions, collect reward. *How should you act?*
- **The common engine:** a single **Bayesian mixture** ξ over a class $\mathcal{M}$ of candidate environments, weighted by a prior w_ν.
  - Solomonoff: ξ (eventually) predicts as well as the truth μ.
  - AIXI: the Bayes-optimal policy π^*_ξ (eventually) acts as well as the optimal policy π^*_μ for true environment μ.

The big picture: one agent, two halves

- **Universal AI** asks: what would an *optimally intelligent* agent look like, given unlimited compute and the weakest possible assumptions?
- It splits into two problems:
  1. **Prediction (Solomonoff):** passively observe a sequence $x_1, x_2, \ldots$ and predict the next symbol. *How should you learn?*
  2. **Action (AIXI):** interact with an unknown environment, take actions, collect reward. *How should you act?*
- **The common engine:** a single **Bayesian mixture** ξ over a class $\mathcal{M}$ of candidate environments, weighted by a prior w_ν.
  - Solomonoff: ξ (eventually) predicts as well as the truth μ.
  - AIXI: the Bayes-optimal policy π^*_ξ (eventually) acts as well as the optimal policy π^*_μ for true environment μ.
- **Universal** choice: $\mathcal{M} =$ all computable environments, prior $w_\nu = 2^{-K(\nu)}$ (K is the "complexity" of the environment). Assumption "$\mu \in \mathcal{M}$" becomes "*the universe is computable*".

The Bayesian mixture ξ

Definition

Class $\mathcal{M} = \{\nu_1, \nu_2, \dots\}$ with prior weights $w_\nu > 0$, $\sum_\nu w_\nu = 1$.

$$\xi(x_{1:t}) := \sum_{\nu \in \mathcal{M}} w_\nu \, \nu(x_{1:t}).$$

The Bayesian mixture ξ

Definition

Class $\mathcal{M} = \{\nu_1, \nu_2, \dots\}$ with prior weights $w_\nu > 0$, $\sum_\nu w_\nu = 1$.

$$\xi(x_{1:t}) := \sum_{\nu \in \mathcal{M}} w_\nu \, \nu(x_{1:t}).$$

- **Key fact (mixture is a predictor):** its one-step prediction is a *posterior-weighted* average of the members' predictions (you prove this today!)

$$\xi(x_t \mid x_{<t}) = \sum_{\nu \in \mathcal{M}} w(\nu \mid x_{<t}) \, \nu(x_t \mid x_{<t}), \qquad w(\nu \mid x_{<t}) = w_\nu \, \frac{\nu(x_{<t})}{\xi(x_{<t})}.$$

The Bayesian mixture ξ

Definition

Class $\mathcal{M} = \{\nu_1, \nu_2, \dots\}$ with prior weights $w_\nu > 0$, $\sum_\nu w_\nu = 1$.

$$\xi(x_{1:t}) := \sum_{\nu \in \mathcal{M}} w_\nu \, \nu(x_{1:t}).$$

- **Key fact (mixture is a predictor):** its one-step prediction is a *posterior-weighted* average of the members' predictions (you prove this today!)

$$\xi(x_t \mid x_{<t}) = \sum_{\nu \in \mathcal{M}} w(\nu \mid x_{<t}) \, \nu(x_t \mid x_{<t}), \qquad w(\nu \mid x_{<t}) = w_\nu \, \frac{\nu(x_{<t})}{\xi(x_{<t})}.$$

- **Posterior update is multiplicative:** $w(\nu \mid x_{1:t}) = w(\nu \mid x_{<t}) \cdot \dfrac{\nu(x_t \mid x_{<t})}{\xi(x_t \mid x_{<t})}$: we reweight by how well ν predicted vs. ξ (Bayes' rule!)

From prediction to action

- Now the agent doesn't just watch: it **acts**, and its actions shape what it sees.

From prediction to action

- Now the agent doesn't just watch: it **acts**, and its actions shape what it sees.
- At each step t: the agent picks action a_t, the environment returns a **percept** $e_t = (o_t, r_t) =$ observation + reward.

$$h = a_1\, e_1\, a_2\, e_2\, a_3\, e_3\, \cdots \qquad \text{(history } æ_{<t}\text{)}$$

From prediction to action

- Now the agent doesn't just watch: it **acts**, and its actions shape what it sees.
- At each step t: the agent picks action a_t, the environment returns a **percept** $e_t = (o_t, r_t) =$ observation + reward.

$$h = a_1\, e_1\, a_2\, e_2\, a_3\, e_3\, \cdots \qquad \text{(history } æ_{<t}\text{)}$$

- Spaces: actions $\mathcal{A}$, observations $\mathcal{O}$, rewards $\mathcal{R} \subset [0,1]$, percepts $\mathcal{E} = \mathcal{O} \times \mathcal{R}$, and (finite) histories $\mathcal{H}^* := (\mathcal{A} \times \mathcal{E})^* := \cup_{t=0}^{\infty} (\mathcal{A} \times \mathcal{E})^t$.

From prediction to action

- Now the agent doesn't just watch: it **acts**, and its actions shape what it sees.
- At each step t: the agent picks action a_t, the environment returns a **percept** $e_t = (o_t, r_t) =$ observation + reward.

$$h = a_1\, e_1\, a_2\, e_2\, a_3\, e_3\, \cdots \qquad \text{(history } æ_{<t}\text{)}$$

- Spaces: actions $\mathcal{A}$, observations $\mathcal{O}$, rewards $\mathcal{R} \subset [0,1]$, percepts $\mathcal{E} = \mathcal{O} \times \mathcal{R}$, and (finite) histories $\mathcal{H}^* := (\mathcal{A} \times \mathcal{E})^* := \cup_{t=0}^{\infty} (\mathcal{A} \times \mathcal{E})^t$.
- **No Markov / MDP assumption!** Both the policy and environment can depend on all prior percepts.

The three players

Policy π

$\pi(a_t \mid æ_{<t}) : \mathcal{H} \to \Delta\mathcal{A}$ is the agent's distribution over the next action given history.

The three players

Policy π

$\pi(a_t \mid æ_{<t}) : \mathcal{H} \to \Delta\mathcal{A}$ is the agent's distribution over the next action given history.

Environment ν

$\nu(e_t \mid æ_{<t}\, a_t) : \mathcal{H} \times \mathcal{A} \to \Delta\mathcal{E}$ is the environment's distribution over the next percept given history and action. True (unknown) environment is μ.

The three players

Policy π

$\pi(a_t \mid æ_{<t}) : \mathcal{H} \to \Delta\mathcal{A}$ is the agent's distribution over the next action given history.

Environment ν

$\nu(e_t \mid æ_{<t}\, a_t) : \mathcal{H} \times \mathcal{A} \to \Delta\mathcal{E}$ is the environment's distribution over the next percept given history and action. True (unknown) environment is μ.

Interaction measure ν^π

Run π inside ν and multiply along the history:

$$\nu^\pi(æ_{t:m} \mid æ_{<t}) := \prod_{k=t}^{m} \pi(a_k \mid æ_{<k})\, \nu(e_k \mid æ_{<k} a_k).$$

Value, optimality, and AIXI

Value function (assuming geometric discount γ)

$$V_\nu^{\pi,m}(æ_{<t}) = (1-\gamma)\,\mathbb{E}_\nu^\pi[G_{t-1:m} \mid æ_{<t}] \text{ with } G_{t-1:m} = \sum_{k=t}^{m} \gamma^{k-t}\, r_k$$

$$= (1-\gamma) \sum_{æ_{t:m}} \nu^\pi(æ_{t:m} \mid æ_{<t}) \sum_{k=t}^{m} \gamma^{k-t}\, r_k$$

Value, optimality, and AIXI

Value function (assuming geometric discount γ)

$$V_\nu^{\pi,m}(æ_{<t}) = (1-\gamma)\,\mathbb{E}_\nu^\pi[G_{t-1:m} \mid æ_{<t}] \text{ with } G_{t-1:m} = \sum_{k=t}^{m} \gamma^{k-t}\, r_k$$

$$= (1-\gamma) \sum_{æ_{t:m}} \nu^\pi(æ_{t:m} \mid æ_{<t}) \sum_{k=t}^{m} \gamma^{k-t}\, r_k$$

The **infinite-horizon value** is defined as the pointwise limit

$$V_\nu^\pi \equiv V_\nu^{\pi,\infty}(æ_{<t}) := \lim_{m\to\infty} V_\nu^{\pi,m}(æ_{<t}).$$

Value, optimality, and AIXI

Value function (assuming geometric discount γ)

$$V_\nu^{\pi,m}(æ_{<t}) = (1-\gamma)\,\mathbb{E}_\nu^\pi[G_{t-1:m} \mid æ_{<t}] \text{ with } G_{t-1:m} = \sum_{k=t}^{m} \gamma^{k-t}\, r_k$$

$$= (1-\gamma) \sum_{æ_{t:m}} \nu^\pi(æ_{t:m} \mid æ_{<t}) \sum_{k=t}^{m} \gamma^{k-t}\, r_k$$

The **infinite-horizon value** is defined as the pointwise limit

$$V_\nu^\pi \equiv V_\nu^{\pi,\infty}(æ_{<t}) := \lim_{m\to\infty} V_\nu^{\pi,m}(æ_{<t}).$$

The **optimal value** is $V_\nu^{*,m}(æ_{<t}) := \sup_\pi V_\nu^{\pi,m}(æ_{<t})$, with $V_\nu^* \equiv V_\nu^{*,\infty}$.

Value, optimality, and AIXI

Value function (assuming geometric discount γ)

$$V_\nu^{\pi,m}(æ_{<t}) = (1-\gamma)\,\mathbb{E}_\nu^\pi[G_{t-1:m} \mid æ_{<t}] \text{ with } G_{t-1:m} = \sum_{k=t}^{m} \gamma^{k-t}\, r_k$$

$$= (1-\gamma) \sum_{æ_{t:m}} \nu^\pi(æ_{t:m} \mid æ_{<t}) \sum_{k=t}^{m} \gamma^{k-t}\, r_k$$

The **infinite-horizon value** is defined as the pointwise limit

$$V_\nu^\pi \equiv V_\nu^{\pi,\infty}(æ_{<t}) := \lim_{m\to\infty} V_\nu^{\pi,m}(æ_{<t}).$$

The **optimal value** is $V_\nu^{*,m}(æ_{<t}) := \sup_\pi V_\nu^{\pi,m}(æ_{<t})$, with $V_\nu^* \equiv V_\nu^{*,\infty}$.

- **Optimal value** $V_\nu^* := \sup_\pi V_\nu^\pi$; an optimal policy π_ν^* satisfies $V_\nu^{\pi_\nu^*} = V_\nu^*$.
- **Bayes-optimal policy** $\pi_\xi^* \in \arg\max_\pi V_\xi^\pi$: acts optimally w.r.t. the *mixture* ξ.

AIXI

π^*_ξ with $\mathcal{M}$ = all ≈computable environments[a] and prior $w_\nu = 2^{-K(\nu)}$. Write ξ_U for ξ with this choice of $\mathcal{M}$ and w_ν, so AIXI $= \pi^*_{\xi_U}$.

[a] lower-semicomputable semi-measures

Intuition: *Act optimally according to environments that you believe to be in, assuming the simplest environments consistent with observations are more likely to be true.*

AIXI

AIXI

π^*_ξ with $\mathcal{M}$ = all ≈computable environments[a] and prior $w_\nu = 2^{-K(\nu)}$. Write ξ_U for ξ with this choice of $\mathcal{M}$ and w_ν, so AIXI $= \pi^*_{\xi_U}$.

[a] lower-semicomputable semi-measures

Intuition: *Act optimally according to environments that you believe to be in, assuming the simplest environments consistent with observations are more likely to be true.*

AIXI, explicitly

max over actions, $\sum$ over percepts, choosing actions that maximise the ξ-weighted discounted return:

$$a_t = \arg\max_{a_t} \lim_{m\to\infty} \sum_{e_t} \max_{a_{t+1}} \sum_{e_{t+1}} \cdots \max_{a_m} \sum_{e_m} \xi(e_{t:m} \mid æ_{<t}\, a_{t:m}) \sum_{k=t}^{m} \gamma^{k-t} r_k$$

Result 1: On-policy value convergence

Theorem

For any $\mu \in \mathcal{M}$ and any policy π,

$$V_\xi^\pi(æ_{<t}) - V_\mu^\pi(æ_{<t}) \to 0 \qquad \mu^\pi\text{-almost surely.}$$

Intuition: *Following policy π, pretending the environment is ξ, eventually will act as if we knew the true environment μ.*

Result 2: AIXI can't be fooled

Theorem (deterministic μ)

If $V^*_\mu(æ_{<t}) > \varepsilon$ along history $æ_{<t}$, then

$$V^*_\xi(æ_{<t}) \;\geq\; w_\mu\,\varepsilon \;>\; 0.$$

Result 2: AIXI can't be fooled

Theorem (deterministic μ)

If $V_\mu^*(æ_{<t}) > \varepsilon$ along history $æ_{<t}$, then

$$V_\xi^*(æ_{<t}) \;\geq\; w_\mu\,\varepsilon \;>\; 0.$$

Intuition: *Whenever good behaviour in μ is achievable, the Bayes-optimal agent secures at least a fraction proportional to the prior weight of μ.*

Result 3: Self-optimizing property (stretch goal)

Theorem

If *some* policy can learn to act optimally in every $\nu \in \mathcal{M}$, then π^*_ξ does too:

$$V^*_\mu(æ_{<t}) - V^{\pi^*_\xi}_\mu(æ_{<t}) \to 0 \quad \mu^\pi\text{-a.s.}$$

Result 3: Self-optimizing property (stretch goal)

Theorem

If *some* policy can learn to act optimally in every $\nu \in \mathcal{M}$, then π^*_ξ does too:

$$V^*_\mu(æ_{<t}) - V^{\pi^*_\xi}_\mu(æ_{<t}) \to 0 \quad \mu^\pi\text{-a.s.}$$

Intuition: *AIXI learns to play as well as the optimal agent that knows the true environment* μ.

Takeaways

- **Solomonoff** solves induction: mix over all computable hypotheses, weighted by simplicity. Total prediction error $\leq K(\mu) \ln 2$.

Takeaways

- **Solomonoff** solves induction: mix over all computable hypotheses, weighted by simplicity. Total prediction error $\leq K(\mu) \ln 2$.
- **AIXI** lifts this to action: the Bayes-optimal policy over the same universal mixture. It learns to predict *and* to act as well as if it knew the world.

Takeaways

- **Solomonoff** solves induction: mix over all computable hypotheses, weighted by simplicity. Total prediction error $\leq K(\mu) \ln 2$.
- **AIXI** lifts this to action: the Bayes-optimal policy over the same universal mixture. It learns to predict *and* to act as well as if it knew the world.
- **The catch:** both are *incomputable* (K is incomputable; $\mathcal{M}$ is intractable to enumerate, the expectimax is intractable). They are gold standards, not computable algorithms, and the prior's $2^{-K(\nu)}$ hides a dependency on the choice of universal Turing machine.

Takeaways

- **Solomonoff** solves induction: mix over all computable hypotheses, weighted by simplicity. Total prediction error $\leq K(\mu) \ln 2$.
- **AIXI** lifts this to action: the Bayes-optimal policy over the same universal mixture. It learns to predict *and* to act as well as if it knew the world.
- **The catch:** both are *incomputable* (K is incomputable; $\mathcal{M}$ is intractable to enumerate, the expectimax is intractable). They are gold standards, not computable algorithms, and the prior's $2^{-K(\nu)}$ hides a dependency on the choice of universal Turing machine.
- **Why care for safety:** a precise, assumption-light model of idealized intelligence providing a lens on what optimal agents do. For example, we can show under what circumstances AIXI would wirehead [RO11],
- **Approximations:** MC-AIXI-CTW [VNH+11] gives an approximation using Context Tree Weighting to replace ξ and Monte-Carlo-Tree-Search to replace the expectimax. safely rather than recklessly explore [CCH20], or self-modify [OR11] as a model of what superintelligent agents would do.

References I

[CCH20] Michael K. Cohen, Elliot Catt, and Marcus Hutter.
Curiosity killed or incapacitated the cat and the asymptotically optimal agent.
CoRR, abs/2006.03357, 2020.
URL: https://arxiv.org/abs/2006.03357.

[OR11] Laurent Orseau and Mark Ring.
Self-modification and mortality in artificial agents.
In *Artificial General Intelligence (AGI 2011)*, volume 6830 of *Lecture Notes in Computer Science*, pages 1–10. Springer, 2011.
URL: https://link.springer.com/chapter/10.1007/978-3-642-22887-2_1, doi:10.1007/978-3-642-22887-2_1.

[RO11] Mark Ring and Laurent Orseau.
Delusion, survival, and intelligent agents.
In *Artificial General Intelligence (AGI 2011)*, volume 6830 of *Lecture Notes in Computer Science*, pages 11–20. Springer, 2011.
URL: https://people.idsia.ch/~ring/AGI-2011/Paper-B.pdf, doi:10.1007/978-3-642-22887-2_2.

References II

[VNH+11] Joel Veness, Kee Siong Ng, Marcus Hutter, William Uther, and David Silver.
A Monte-Carlo AIXI approximation.
Journal of Artificial Intelligence Research, 40:95–142, 2011.
URL: https://arxiv.org/abs/0909.0801, doi:10.1613/jair.3125.